

```
#In case google denies access for downloading the file, you can try by installing the l
pip install -U --no-cache-dir gdown --pre

  File "<ipython-input-2-336916444db0>", line 2
    pip install -U --no-cache-dir gdown --pre
    ^
SyntaxError: invalid syntax
```

SEARCH STACK OVERFLOW

```
#Download the dataset
#Shared file link https://drive.google.com/file/d/1WfdEh2CyoV-cSNFyncKvLYhmD0rv2fFN/view
#Alternatively you can download to your local computer an upload back to colab

!gdown --id '1WfdEh2CyoV-cSNFyncKvLYhmD0rv2fFN' --output acid_sub

  Downloading...
  From: https://drive.google.com/uc?id=1WfdEh2CyoV-cSNFyncKvLYhmD0rv2fFN
  To: /content/acid_sub
  1.25GB [00:10, 124MB/s]
```

▼ Object Detection tutorial

For this tutorial, we will be finetuning a pre-trained [Faster R-CNN](#) model in the ACID dataset with 3 classes. It contains 3929 images with 8000+ instances of construction equipment, and we will use it to illustrate how to use the new features in torchvision in order to train an object detection model on a custom dataset.

First, we need to install `pycocotools`. This library will be used for computing the evaluation metrics following the COCO metric for intersection over union.

```
%%shell

pip install cython
# Install pycocotools, the version by default in Colab
# has a bug fixed in https://github.com/cocodataset/cocoapi/pull/354
pip install -U 'git+https://github.com/cocodataset/cocoapi.git#subdirectory=PythonAPI'
```

```
Requirement already satisfied: cython in /usr/local/lib/python3.7/dist-packages ((

Collecting git+https://github.com/cocodataset/cocoapi.git#subdirectory=PythonAPI
  Cloning https://github.com/cocodataset/cocoapi.git to /tmp/pip-req-build-_0c2mak
    Running command git clone -q https://github.com/cocodataset/cocoapi.git /tmp/pip-req-build-_0c2mak
Requirement already satisfied, skipping upgrade: setuptools>=18.0 in /usr/local/lib/python3.7/dist-packages/setuptools-44.1.1-py3.7.egg
Requirement already satisfied, skipping upgrade: cython>=0.27.3 in /usr/local/lib/python3.7/dist-packages/cython-0.29.21-py3.7.egg
Requirement already satisfied, skipping upgrade: matplotlib>=2.1.0 in /usr/local/lib/python3.7/dist-packages/matplotlib-3.3.2-py3.7.egg
Requirement already satisfied, skipping upgrade: cycler>=0.10 in /usr/local/lib/python3.7/dist-packages/cycler-0.10.0-py3.7.egg
Requirement already satisfied, skipping upgrade: python-dateutil>=2.1 in /usr/local/lib/python3.7/dist-packages/python_dateutil-2.8.1-py3.7.egg
Requirement already satisfied, skipping upgrade: kiwisolver>=1.0.1 in /usr/local/lib/python3.7/dist-packages/kiwisolver-1.3.1-py3.7.egg
Requirement already satisfied, skipping upgrade: numpy>=1.11 in /usr/local/lib/python3.7/dist-packages/numpy-1.19.2-py3.7.egg
Requirement already satisfied, skipping upgrade: pyparsing!=2.0.4,>2.1.2,<2.1.6,!=2.1.4 in /usr/local/lib/python3.7/dist-packages/pyparsing-2.4.7-py3.7.egg
Requirement already satisfied, skipping upgrade: six in /usr/local/lib/python3.7/dist-packages/six-1.15.0-py3.7.egg
Building wheels for collected packages: pycocotools
  Building wheel for pycocotools (setup.py) ... done
    Created wheel for pycocotools: filename=pycocotools-2.0-cp37-cp37m-linux_x86_64.
```

```
Stored in directory: /tmp/pip-ephem-wheel-cache-vpdmakm0/wheels/90/51/41/646daf4
Successfully built pycocotools
Installing collected packages: pycocotools
  Found existing installation: pycocotools 2.0.2
    Uninstalling pycocotools-2.0.2:
      Successfully uninstalled pycocotools-2.0.2
Successfully installed pycocotools-2.0
```

▼ Defining the Dataset

The [torchvision reference scripts for training object detection, instance segmentation and person keypoint detection](#) allows for easily supporting adding new custom datasets. The dataset should inherit from the standard `torch.utils.data.Dataset` class, and implement `__len__` and `__getitem__`.

The only specificity that we require is that the dataset `__getitem__` should return:

- `image`: a PIL Image of size (`H`, `W`)
- `target`: a dict containing the following fields
 - `boxes` (`FloatTensor[N, 4]`): the coordinates of the `N` bounding boxes in `[x0, y0, x1, y1]` format, ranging from 0 to `w` and 0 to `H`
 - `labels` (`Int64Tensor[N]`): the label for each bounding box
 - `image_id` (`Int64Tensor[1]`): an image identifier. It should be unique between all the images in the dataset, and is used during evaluation
 - `area` (`Tensor[N]`): The area of the bounding box. This is used during evaluation with the COCO metric, to separate the metric scores between small, medium and large boxes.
 - `iscrowd` (`UInt8Tensor[N]`): instances with `iscrowd=True` will be ignored during evaluation.
 - (optionally) `masks` (`UInt8Tensor[N, H, W]`): The segmentation masks for each one of the objects
 - (optionally) `keypoints` (`FloatTensor[N, K, 3]`): For each one of the `N` objects, it contains the `K` keypoints in `[x, y, visibility]` format, defining the object. `visibility=0` means that the keypoint is not visible. Note that for data augmentation, the notion of flipping a keypoint is dependent on the data representation, and you should probably adapt `references/detection/transforms.py` for your new keypoint representation

If your model returns the above methods, they will make it work for both training and evaluation, and will use the evaluation scripts from pycocotools.

Additionally, if you want to use aspect ratio grouping during training (so that each batch only contains images with similar aspect ratio), then it is recommended to also implement a `get_height_and_width` method, which returns the height and the width of the image. If this method is not provided, we query all elements of the dataset via `__getitem__`, which loads the image in memory and is slower than if a custom method is provided.

▼ Writing a custom dataset for ACID

Let's write a dataset for the ACID dataset.

First, let's extract the data, present in a zip file

```
!unzip acid_sub
-----, -----
inflating: 3classes/08686.jpg
inflating: 3classes/08688.jpg
inflating: 3classes/08689.jpg
inflating: 3classes/08690.jpg
inflating: 3classes/08695.jpg
inflating: 3classes/08696.jpg
inflating: 3classes/08699.jpg
inflating: 3classes/08700.jpg
inflating: 3classes/08701.jpg
inflating: 3classes/08703.jpg
inflating: 3classes/08705.jpg
inflating: 3classes/08707.jpg
inflating: 3classes/08713.jpg
inflating: 3classes/08714.jpg
inflating: 3classes/08717.jpg

inflating: 3classes/08720.jpg
inflating: 3classes/08722.jpg
inflating: 3classes/08723.jpg
inflating: 3classes/08725.jpg
inflating: 3classes/08726.jpg
inflating: 3classes/08727.jpg
inflating: 3classes/08728.jpg
inflating: 3classes/08731.jpg
inflating: 3classes/08733.jpg
inflating: 3classes/08737.jpg
inflating: 3classes/08745.jpg
inflating: 3classes/08746.jpg
inflating: 3classes/08758.jpg
inflating: 3classes/08759.jpg
inflating: 3classes/08760.jpg
inflating: 3classes/08763.jpg
inflating: 3classes/08768.jpg
inflating: 3classes/08770.jpg
inflating: 3classes/08772.jpg
inflating: 3classes/08773.jpg
inflating: 3classes/08777.jpg
inflating: 3classes/08778.jpg
inflating: 3classes/08788.jpg
inflating: 3classes/08791.jpg
inflating: 3classes/08793.jpg
inflating: 3classes/08795.jpg
inflating: 3classes/08796.jpg
inflating: 3classes/08797.jpg
inflating: 3classes/08799.jpg
inflating: 3classes/08805.jpg
inflating: 3classes/08810.jpg
inflating: 3classes/08811.jpg
inflating: 3classes/08812.jpg
inflating: 3classes/08816.jpg
inflating: 3classes/08818.jpg
inflating: 3classes/08828.jpg
inflating: 3classes/08831.jpg
inflating: 3classes/08832.jpg
inflating: 3classes/08833.jpg
inflating: 3classes/08836.jpg
inflating: 3classes/08840.jpg
inflating: 3classes/08841.jpg
inflating: 3classes/08842.jpg
```

```
inflating: 3classes/0001.jpg
inflating: 3classes/08846.jpg
```

Let's have a look at the dataset and how it is layed down.

The data is structured as follows

```
3classes/
    00001.jpg
    00002.jpg
    ....
    ....
    3classes.csv
    3classes.json
```

Here is one example of an image in the dataset

```
from PIL import Image
import os
path = '/content/3classes/'
Image.open(os.path.join(path, '10031.jpg'))
#Image.open(path)
```



So each image has a corresponding bounding box. Let's write a `torch.utils.data.Dataset` class for this dataset.



```

import os
import numpy as np
import torch
import torch.utils.data
from PIL import Image
from pycocotools.coco import COCO

class ACID_dataset(torch.utils.data.Dataset):
    def __init__(self, root, annotation, transforms=None):
        self.root = root
        self.transforms = transforms
        self.coco = COCO(annotation)
        self.ids = list(sorted(self.coco.imgs.keys()))

    def __getitem__(self, idx):
        # Own coco file
        coco = self.coco
        # Image ID
        img_id = self.ids[idx]
        # List: get annotation id from coco
        ann_ids = coco.getAnnIds(imgIds=img_id)

        # Dictionary: target coco_annotation file for an image
        coco_annotation = coco.loadAnns(ann_ids)

        # path for input image
        path = coco.loadImgs(img_id)[0]['file_name']
        # open the input image
        img = Image.open(os.path.join(self.root, path))

        # get bounding box coordinates for each object
        num_objs = len(ann_ids)
        num_objs = len(coco_annotation)

        boxes = []
        cat_ids = []
        for i in range(num_objs):
            xmin = coco_annotation[i]['bbox'][0]
            ymin = coco_annotation[i]['bbox'][1]
            xmax = xmin + coco_annotation[i]['bbox'][2]
            ymax = ymin + coco_annotation[i]['bbox'][3]
            boxes.append([xmin, ymin, xmax, ymax])
            cat_id = coco_annotation[i]["category_id"]
            cat_ids.append(cat_id)
        boxes = torch.as_tensor(boxes, dtype=torch.float32)

        return img, boxes, cat_ids

```

```

labels = torch.as_tensor(cat_ids, dtype=torch.int64)

image_id = torch.tensor([idx])
area = (boxes[:, 3] - boxes[:, 1]) * (boxes[:, 2] - boxes[:, 0])
# suppose all instances are not crowd
iscrowd = torch.zeros((num_objs,), dtype=torch.int64)

target = {}
target["boxes"] = boxes
target["labels"] = labels
target["image_id"] = image_id
target["area"] = area
target["iscrowd"] = iscrowd

if self.transforms is not None:
    img, target = self.transforms(img, target)

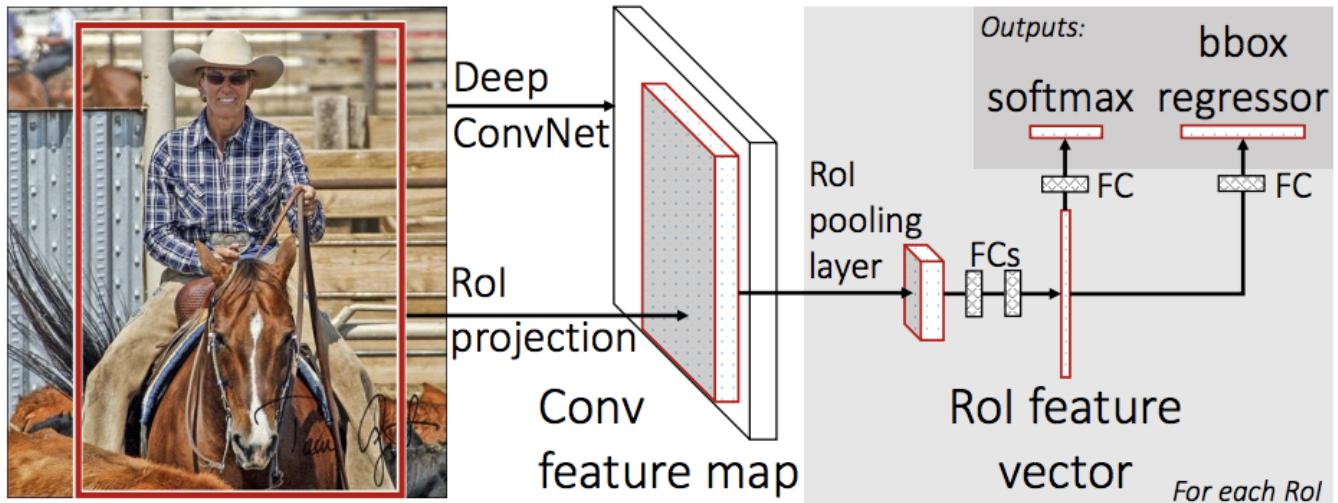
return img, target

def __len__(self):
    return len(self.ids)

```

▼ Defining your model

In this tutorial, we will be using [Faster R-CNN](#). Faster R-CNN is a model that predicts both bounding boxes and class scores for potential objects in the image.



There are two common situations where one might want to modify one of the available models in torchvision modelzoo. The first is when we want to start from a pre-trained model, and just finetune the last layer. The other is when we want to replace the backbone of the model with a different one (for faster predictions, for example).

Let's go see how we would do one or another in the following sections.

1 - Finetuning from a pretrained model

Let's suppose that you want to start from a model pre-trained on COCO and want to finetune it for your particular classes. Here is a possible way of doing it:

```

import torchvision
from torchvision.models.detection.faster_rcnn import FastRCNNPredictor

# load a model pre-trained pre-trained on COCO
model = torchvision.models.detection.fasterrcnn_resnet50_fpn(pretrained=True)

# replace the classifier with a new one, that has
# num_classes which is user-defined
num_classes = 2 # 1 class (person) + background
# get number of input features for the classifier
in_features = model.roi_heads.box_predictor.cls_score.in_features
# replace the pre-trained head with a new one
model.roi_heads.box_predictor = FastRCNNPredictor(in_features, num_classes)

```

2 - Modifying the model to add a different backbone

Another common situation arises when the user wants to replace the backbone of a detection model with a different one. For example, the current default backbone (ResNet-50) might be too big for some applications, and smaller models might be necessary.

Here is how we would go into leveraging the functions provided by torchvision to modify a backbone.

```

import torchvision
from torchvision.models.detection import FasterRCNN
from torchvision.models.detection.rpn import AnchorGenerator

# load a pre-trained model for classification and return
# only the features
backbone = torchvision.models.mobilenet_v2(pretrained=True).features
# FasterRCNN needs to know the number of
# output channels in a backbone. For mobilenet_v2, it's 1280
# so we need to add it here
backbone.out_channels = 1280

# let's make the RPN generate 5 x 3 anchors per spatial
# location, with 5 different sizes and 3 different aspect
# ratios. We have a Tuple[Tuple[int]] because each feature
# map could potentially have different sizes and
# aspect ratios
anchor_generator = AnchorGenerator(sizes=((32, 64, 128, 256, 512),),
                                    aspect_ratios=((0.5, 1.0, 2.0),))

# let's define what are the feature maps that we will
# use to perform the region of interest cropping, as well as
# the size of the crop after rescaling.
# if your backbone returns a Tensor, featmap_names is expected to
# be [0]. More generally, the backbone should return an
# OrderedDict[Tensor], and in featmap_names you can choose which
# feature maps to use.

```

```

roi_pooler = torchvision.ops.MultiScaleRoIAlign(featmap_names=[0],
                                                output_size=7,
                                                sampling_ratio=2)

# put the pieces together inside a FasterRCNN model
model = FasterRCNN(backbone,
                     num_classes=2,
                     rpn_anchor_generator=anchor_generator,
                     box_roi_pool=roi_pooler)

```

An object_detection model for Construction Equipment Detection

In our case, we want to fine-tune from a pre-trained model, given that our dataset is very small. So we will be following approach number 1.

```

import torchvision
from torchvision.models.detection.faster_rcnn import FastRCNNPredictor

def get_object_detection_model(num_classes):
    # load an object_detection model pre-trained on COCO
    model = torchvision.models.detection.fasterrcnn_resnet50_fpn(pretrained=True)
    # get the number of input features for the classifier
    in_features = model.roi_heads.box_predictor.cls_score.in_features
    # replace the pre-trained head with a new one
    model.roi_heads.box_predictor = FastRCNNPredictor(in_features, num_classes)

    return model

```

That's it, this will make model be ready to be trained and evaluated on our custom dataset.

▼ Training and evaluation functions

In `references/detection/`, we have a number of helper functions to simplify training and evaluating detection models. Here, we will use `references/detection/engine.py`, `references/detection/utils.py` and `references/detection/transforms.py`.

Let's copy those files (and their dependencies) in here so that they are available in the notebook

```

%%shell

# Download TorchVision repo to use some files from
# references/detection
git clone https://github.com/pytorch/vision.git
cd vision
git checkout v0.3.0

cp references/detection/utils.py ../
cp references/detection/transforms.py ../
cp references/detection/coco_eval.py ../
cp references/detection/engine.py ../

```

```
cp references/detection/coco_utils.py ...
```

```
Cloning into 'vision'...
remote: Enumerating objects: 26939, done.
remote: Counting objects: 100% (4548/4548), done.
remote: Compressing objects: 100% (1159/1159), done.
remote: Total 26939 (delta 3508), reused 4261 (delta 3305), pack-reused 22391
Receiving objects: 100% (26939/26939), 35.03 MiB | 34.46 MiB/s, done.
Resolving deltas: 100% (20082/20082), done.
Note: checking out 'v0.3.0'.
```

You are in 'detached HEAD' state. You can look around, make experimental changes and commit them, and you can discard any commits you make in this state without impacting any branches by performing another checkout.

If you want to create a new branch to retain commits you create, you may do so (now or later) by using -b with the checkout command again. Example:

```
git checkout -b <new-branch-name>
```

```
HEAD is now at be376084 version check against PyTorch's CUDA version
```

Let's write some helper functions for data augmentation / transformation, which leverages the functions in `references/detection` that we have just copied:

```
from engine import train_one_epoch, evaluate
import utils
import transforms as T

def get_transform(train):
    transforms = []
    # converts the image, a PIL image, into a PyTorch Tensor
    transforms.append(T.ToTensor())
    if train:
        # during training, randomly flip the training images
        # and ground-truth for data augmentation
        transforms.append(T.RandomHorizontalFlip(0.5))
    return T.Compose(transforms)
```

Note that we do not need to add a mean/std normalization nor image rescaling in the data transforms, as those are handled internally by the Faster R-CNN model.

▼ Putting everything together

We now have the dataset class, the models and the data transforms. Let's instantiate them

```
# use our dataset and defined transformations
annotation_root = '/content/3classes/3classes.json'
file_root = '/content/3classes'

dataset = ACID_dataset(root = file_root, annotation = annotation_root, transforms = get_
dataset_val = ACID_dataset(root = file_root, annotation = annotation_root, transforms = g
```

```
# split the dataset in train and test set [use 20% for testing]
torch.manual_seed(1)
indices = torch.randperm(len(dataset)).tolist()
print(len(indices))
val_size = int(0.2*len(indices))
print(val_size)
dataset = torch.utils.data.Subset(dataset, indices[:-val_size])
dataset_val = torch.utils.data.Subset(dataset_val, indices[-val_size:])

# define training and validation data loaders
data_loader = torch.utils.data.DataLoader(
    dataset, batch_size=2, shuffle=True, num_workers=4,
    collate_fn=utils.collate_fn)

data_loader_val = torch.utils.data.DataLoader(
    dataset_val, batch_size=1, shuffle=False, num_workers=4,
    collate_fn=utils.collate_fn)

loading annotations into memory...
Done (t=0.03s)
creating index...
index created!
loading annotations into memory...
Done (t=0.02s)
creating index...
index created!
3929
785
/usr/local/lib/python3.7/dist-packages/torch/utils/data/dataloader.py:477: UserWarning
  cpuset_checked))
```

Now let's instantiate the model and the optimizer

```
device = torch.device('cuda') if torch.cuda.is_available() else torch.device('cpu')

# our dataset has 4 classes only - background and 3 types of equipment
num_classes = 4

# get the model using our helper function
model = get_object_detection_model(num_classes)
# move model to the right device
model.to(device)

# construct an optimizer
params = [p for p in model.parameters() if p.requires_grad]
optimizer = torch.optim.SGD(params, lr=0.005,
                            momentum=0.9, weight_decay=0.0005)

# and a learning rate scheduler which decreases the learning rate by
# 10x every 3 epochs
lr_scheduler = torch.optim.lr_scheduler.StepLR(optimizer,
                                                step_size=3,
                                                gamma=0.1)

#Let us visualize the model architecture
print(model)
```

Downloading: "https://download.pytorch.org/models/fasterrcnn_resnet50_fpn_coco-258
100% 160M/160M [00:02<00:00, 66.8MB/s]

```
FasterRCNN(
    (transform): GeneralizedRCNNTransform(
        Normalize(mean=[0.485, 0.456, 0.406], std=[0.229, 0.224, 0.225])
        Resize(min_size=(800,), max_size=1333, mode='bilinear')
    )
    (backbone): BackboneWithFPN(
        (body): IntermediateLayerGetter(
            (conv1): Conv2d(3, 64, kernel_size=(7, 7), stride=(2, 2), padding=(3, 3), bias=False)
            (bn1): FrozenBatchNorm2d(64, eps=0.0)
            (relu): ReLU(inplace=True)
            (maxpool): MaxPool2d(kernel_size=3, stride=2, padding=1, dilation=1, ceil_mode=True)
            (layer1): Sequential(
                (0): Bottleneck(
                    (conv1): Conv2d(64, 64, kernel_size=(1, 1), stride=(1, 1), bias=False)
                    (bn1): FrozenBatchNorm2d(64, eps=0.0)
                    (conv2): Conv2d(64, 64, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)
                    (bn2): FrozenBatchNorm2d(64, eps=0.0)
                    (conv3): Conv2d(64, 256, kernel_size=(1, 1), stride=(1, 1), bias=False)
                    (bn3): FrozenBatchNorm2d(256, eps=0.0)
                    (relu): ReLU(inplace=True)
                    (downsample): Sequential(
                        (0): Conv2d(64, 256, kernel_size=(1, 1), stride=(1, 1), bias=False)
                        (1): FrozenBatchNorm2d(256, eps=0.0)
                    )
                )
            )
            (1): Bottleneck(
                (conv1): Conv2d(256, 64, kernel_size=(1, 1), stride=(1, 1), bias=False)
                (bn1): FrozenBatchNorm2d(64, eps=0.0)
                (conv2): Conv2d(64, 64, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)
                (bn2): FrozenBatchNorm2d(64, eps=0.0)
                (conv3): Conv2d(64, 256, kernel_size=(1, 1), stride=(1, 1), bias=False)
                (bn3): FrozenBatchNorm2d(256, eps=0.0)
                (relu): ReLU(inplace=True)
            )
            (2): Bottleneck(
                (conv1): Conv2d(256, 64, kernel_size=(1, 1), stride=(1, 1), bias=False)
                (bn1): FrozenBatchNorm2d(64, eps=0.0)
                (conv2): Conv2d(64, 64, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)
                (bn2): FrozenBatchNorm2d(64, eps=0.0)
                (conv3): Conv2d(64, 256, kernel_size=(1, 1), stride=(1, 1), bias=False)
                (bn3): FrozenBatchNorm2d(256, eps=0.0)
                (relu): ReLU(inplace=True)
            )
        )
        (layer2): Sequential(
            (0): Bottleneck(
                (conv1): Conv2d(256, 128, kernel_size=(1, 1), stride=(1, 1), bias=False)
                (bn1): FrozenBatchNorm2d(128, eps=0.0)
                (conv2): Conv2d(128, 128, kernel_size=(3, 3), stride=(2, 2), padding=(1, 1), bias=False)
                (bn2): FrozenBatchNorm2d(128, eps=0.0)
                (conv3): Conv2d(128, 512, kernel_size=(1, 1), stride=(1, 1), bias=False)
                (bn3): FrozenBatchNorm2d(512, eps=0.0)
                (relu): ReLU(inplace=True)
                (downsample): Sequential(
                    (0): Conv2d(256, 512, kernel_size=(1, 1), stride=(2, 2), bias=False)
                    (1): FrozenBatchNorm2d(512, eps=0.0)
                )
            )
            (1): Bottleneck(
                (conv1): Conv2d(512, 128, kernel_size=(1, 1), stride=(1, 1), bias=False)
                (bn1): FrozenBatchNorm2d(128, eps=0.0)
            )
        )
    )
)
```

```
r08521609_lab13.ipynb - Colaboratory
  (conv1): Conv2d(128, 128, kernel_size=(3, 3), stride=(1, 1), padding=(1,
  (bn2): FrozenBatchNorm2d(128, eps=0.0)
  (conv3): Conv2d(128, 512, kernel_size=(1, 1), stride=(1, 1), bias=False)
  (bn3): FrozenBatchNorm2d(512, eps=0.0)
  (relu): ReLU(inplace=True)
)
(2): Bottleneck(
  (conv1): Conv2d(512, 128, kernel_size=(1, 1), stride=(1, 1), bias=False)
  (bn1): FrozenBatchNorm2d(128, eps=0.0)
  (conv2): Conv2d(128, 128, kernel_size=(3, 3), stride=(1, 1), padding=(1,
  (bn2): FrozenBatchNorm2d(128, eps=0.0)
  (conv3): Conv2d(128, 512, kernel_size=(1, 1), stride=(1, 1), bias=False)
  (bn3): FrozenBatchNorm2d(512, eps=0.0)
  (relu): ReLU(inplace=True)
)
(3): Bottleneck(
  (conv1): Conv2d(512, 128, kernel_size=(1, 1), stride=(1, 1), bias=False)
  (bn1): FrozenBatchNorm2d(128, eps=0.0)
  (conv2): Conv2d(128, 128, kernel_size=(3, 3), stride=(1, 1), padding=(1,
  (bn2): FrozenBatchNorm2d(128, eps=0.0)
  (conv3): Conv2d(128, 512, kernel_size=(1, 1), stride=(1, 1), bias=False)
  (bn3): FrozenBatchNorm2d(512, eps=0.0)
  (relu): ReLU(inplace=True)
)
)
(layer3): Sequential(
  (0): Bottleneck(
    (conv1): Conv2d(512, 256, kernel_size=(1, 1), stride=(1, 1), bias=False)
    (bn1): FrozenBatchNorm2d(256, eps=0.0)
    (conv2): Conv2d(256, 256, kernel_size=(3, 3), stride=(2, 2), padding=(1,
    (bn2): FrozenBatchNorm2d(256, eps=0.0)
    (conv3): Conv2d(256, 1024, kernel_size=(1, 1), stride=(1, 1), bias=False)
    (bn3): FrozenBatchNorm2d(1024, eps=0.0)
    (relu): ReLU(inplace=True)
    (downsample): Sequential(
      (0): Conv2d(512, 1024, kernel_size=(1, 1), stride=(2, 2), bias=False)
      (1): FrozenBatchNorm2d(1024, eps=0.0)
    )
  )
  (1): Bottleneck(
    (conv1): Conv2d(1024, 256, kernel_size=(1, 1), stride=(1, 1), bias=False)
    (bn1): FrozenBatchNorm2d(256, eps=0.0)
    (conv2): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1,
    (bn2): FrozenBatchNorm2d(256, eps=0.0)
    (conv3): Conv2d(256, 1024, kernel_size=(1, 1), stride=(1, 1), bias=False)
    (bn3): FrozenBatchNorm2d(1024, eps=0.0)
    (relu): ReLU(inplace=True)
  )
  (2): Bottleneck(
    (conv1): Conv2d(1024, 256, kernel_size=(1, 1), stride=(1, 1), bias=False)
    (bn1): FrozenBatchNorm2d(256, eps=0.0)
    (conv2): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1,
    (bn2): FrozenBatchNorm2d(256, eps=0.0)
    (conv3): Conv2d(256, 1024, kernel_size=(1, 1), stride=(1, 1), bias=False)
    (bn3): FrozenBatchNorm2d(1024, eps=0.0)
    (relu): ReLU(inplace=True)
  )
  (3): Bottleneck(
    (conv1): Conv2d(1024, 256, kernel_size=(1, 1), stride=(1, 1), bias=False)
    (bn1): FrozenBatchNorm2d(256, eps=0.0)
    (conv2): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1,
    (bn2): FrozenBatchNorm2d(256, eps=0.0)
    (conv3): Conv2d(256, 1024, kernel_size=(1, 1), stride=(1, 1), bias=False)
    (bn3): FrozenBatchNorm2d(1024, eps=0.0)
    (relu): ReLU(inplace=True)
  )
)
```

```
(relu): ReLU(inplace=True)
)
(4): Bottleneck(
    (conv1): Conv2d(1024, 256, kernel_size=(1, 1), stride=(1, 1), bias=False)
    (bn1): FrozenBatchNorm2d(256, eps=0.0)
    (conv2): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1,
        (bn2): FrozenBatchNorm2d(256, eps=0.0)
    (conv3): Conv2d(256, 1024, kernel_size=(1, 1), stride=(1, 1), bias=False)
    (bn3): FrozenBatchNorm2d(1024, eps=0.0)
    (relu): ReLU(inplace=True)
)
(5): Bottleneck(
    (conv1): Conv2d(1024, 256, kernel_size=(1, 1), stride=(1, 1), bias=False)
    (bn1): FrozenBatchNorm2d(256, eps=0.0)
    (conv2): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1,
        (bn2): FrozenBatchNorm2d(256, eps=0.0)
    (conv3): Conv2d(256, 1024, kernel_size=(1, 1), stride=(1, 1), bias=False)
    (bn3): FrozenBatchNorm2d(1024, eps=0.0)
    (relu): ReLU(inplace=True)
)
)
(layer4): Sequential(
    (0): Bottleneck(
        (conv1): Conv2d(1024, 512, kernel_size=(1, 1), stride=(1, 1), bias=False)
        (bn1): FrozenBatchNorm2d(512, eps=0.0)
        (conv2): Conv2d(512, 512, kernel_size=(3, 3), stride=(2, 2), padding=(1,
            (bn2): FrozenBatchNorm2d(512, eps=0.0)
        (conv3): Conv2d(512, 2048, kernel_size=(1, 1), stride=(1, 1), bias=False)
        (bn3): FrozenBatchNorm2d(2048, eps=0.0)
        (relu): ReLU(inplace=True)
        (downsample): Sequential(
            (0): Conv2d(1024, 2048, kernel_size=(1, 1), stride=(2, 2), bias=False)
            (1): FrozenBatchNorm2d(2048, eps=0.0)
        )
    )
    (1): Bottleneck(
        (conv1): Conv2d(2048, 512, kernel_size=(1, 1), stride=(1, 1), bias=False)
        (bn1): FrozenBatchNorm2d(512, eps=0.0)
        (conv2): Conv2d(512, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1,
            (bn2): FrozenBatchNorm2d(512, eps=0.0)
        (conv3): Conv2d(512, 2048, kernel_size=(1, 1), stride=(1, 1), bias=False)
        (bn3): FrozenBatchNorm2d(2048, eps=0.0)
        (relu): ReLU(inplace=True)
    )
    (2): Bottleneck(
        (conv1): Conv2d(2048, 512, kernel_size=(1, 1), stride=(1, 1), bias=False)
        (bn1): FrozenBatchNorm2d(512, eps=0.0)
        (conv2): Conv2d(512, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1,
            (bn2): FrozenBatchNorm2d(512, eps=0.0)
        (conv3): Conv2d(512, 2048, kernel_size=(1, 1), stride=(1, 1), bias=False)
        (bn3): FrozenBatchNorm2d(2048, eps=0.0)
        (relu): ReLU(inplace=True)
    )
)
)
(fpn): FeaturePyramidNetwork(
    (inner_blocks): ModuleList(
        (0): Conv2d(256, 256, kernel_size=(1, 1), stride=(1, 1))
        (1): Conv2d(512, 256, kernel_size=(1, 1), stride=(1, 1))
        (2): Conv2d(1024, 256, kernel_size=(1, 1), stride=(1, 1))
        (3): Conv2d(2048, 256, kernel_size=(1, 1), stride=(1, 1))
    )
    (layer_blocks): ModuleList(
        (0): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
        (1): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
        (2): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
        (3): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
        (4): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
    )
)
```

```
(2): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
      (3): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
    )
    (extra_blocks): LastLevelMaxPool()
  )
)
(rpn): RegionProposalNetwork(
  (anchor_generator): AnchorGenerator()
  (head): RPNHead(
    (conv): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
    (cls_logits): Conv2d(256, 3, kernel_size=(1, 1), stride=(1, 1))
    (bbox_pred): Conv2d(256, 12, kernel_size=(1, 1), stride=(1, 1))
  )
)
(roi_heads): RoIHeads(
  (box_roi_pool): MultiScaleRoIAlign(featmap_names=['0', '1', '2', '3'], output_
  (box_head): TwoMLPHead(
    (fc6): Linear(in_features=12544, out_features=1024, bias=True)
    (fc7): Linear(in_features=1024, out_features=1024, bias=True)
  )
  (box_predictor): FastRCNNPredictor(
    (cls_score): Linear(in_features=1024, out_features=4, bias=True)
    (bbox_pred): Linear(in_features=1024, out_features=16, bias=True)
  )
)
)
)
```

And now let's train the model for 10 epochs, evaluating at the end of every epoch.

```
! ls /content/3classes

# let's train it for 10 epochs
num_epochs = 1

for epoch in range(num_epochs):
    # train for one epoch, printing every 10 iterations
    train_one_epoch(model, optimizer, data_loader, device, epoch, print_freq=10)
    # update the learning rate
    lr_scheduler.step()
    # evaluate on the test dataset
    evaluate(model, data_loader_val, device=device)

00001.jpg  01527.jpg  02935.jpg  04431.jpg  05853.jpg  07294.jpg  08689.jpg
00002.jpg  01529.jpg  02940.jpg  04433.jpg  05854.jpg  07296.jpg  08690.jpg
00010.jpg  01531.jpg  02941.jpg  04435.jpg  05859.jpg  07297.jpg  08695.jpg
00011.jpg  01534.jpg  02946.jpg  04437.jpg  05861.jpg  07300.jpg  08696.jpg
00012.jpg  01536.jpg  02948.jpg  04438.jpg  05862.jpg  07306.jpg  08699.jpg
00013.jpg  01537.jpg  02949.jpg  04439.jpg  05869.jpg  07307.jpg  08700.jpg
00018.jpg  01538.jpg  02950.jpg  04440.jpg  05870.jpg  07308.jpg  08701.jpg
00021.jpg  01539.jpg  02951.jpg  04444.jpg  05871.jpg  07309.jpg  08703.jpg
00022.jpg  01542.jpg  02952.jpg  04448.jpg  05876.jpg  07313.jpg  08705.jpg
```

00025.jpg	01544.jpg	02959.jpg	04450.jpg	05881.jpg	07315.jpg	08707.jpg
00028.jpg	01546.jpg	02962.jpg	04451.jpg	05883.jpg	07318.jpg	08713.jpg
00030.jpg	01549.jpg	02964.jpg	04452.jpg	05890.jpg	07319.jpg	08714.jpg
00031.jpg	01551.jpg	02972.jpg	04458.jpg	05891.jpg	07321.jpg	08717.jpg
00034.jpg	01556.jpg	02979.jpg	04460.jpg	05893.jpg	07323.jpg	08720.jpg
00037.jpg	01561.jpg	02983.jpg	04461.jpg	05894.jpg	07325.jpg	08722.jpg
00038.jpg	01562.jpg	02984.jpg	04465.jpg	05896.jpg	07326.jpg	08723.jpg
00041.jpg	01564.jpg	02986.jpg	04468.jpg	05897.jpg	07329.jpg	08725.jpg
00042.jpg	01565.jpg	02989.jpg	04471.jpg	05902.jpg	07331.jpg	08726.jpg
00043.jpg	01566.jpg	02992.jpg	04475.jpg	05903.jpg	07332.jpg	08727.jpg
00047.jpg	01568.jpg	02996.jpg	04477.jpg	05904.jpg	07333.jpg	08728.jpg
00051.jpg	01570.jpg	02997.jpg	04482.jpg	05906.jpg	07335.jpg	08731.jpg
00053.jpg	01574.jpg	03000.jpg	04484.jpg	05907.jpg	07337.jpg	08733.jpg
00059.jpg	01576.jpg	03003.jpg	04485.jpg	05911.jpg	07340.jpg	08737.jpg
00066.jpg	01577.jpg	03006.jpg	04486.jpg	05918.jpg	07345.jpg	08745.jpg
00076.jpg	01579.jpg	03008.jpg	04487.jpg	05921.jpg	07346.jpg	08746.jpg
00080.jpg	01580.jpg	03010.jpg	04491.jpg	05922.jpg	07349.jpg	08758.jpg
00081.jpg	01583.jpg	03011.jpg	04493.jpg	05926.jpg	07352.jpg	08759.jpg
00082.jpg	01585.jpg	03014.jpg	04494.jpg	05933.jpg	07355.jpg	08760.jpg
00085.jpg	01587.jpg	03015.jpg	04496.jpg	05937.jpg	07358.jpg	08763.jpg
00088.jpg	01588.jpg	03018.jpg	04497.jpg	05944.jpg	07361.jpg	08768.jpg
00090.jpg	01590.jpg	03021.jpg	04500.jpg	05945.jpg	07365.jpg	08770.jpg
00092.jpg	01591.jpg	03022.jpg	04503.jpg	05950.jpg	07366.jpg	08772.jpg
00093.jpg	01592.jpg	03025.jpg	04504.jpg	05954.jpg	07367.jpg	08773.jpg
00094.jpg	01596.jpg	03026.jpg	04506.jpg	05963.jpg	07368.jpg	08777.jpg
00095.jpg	01597.jpg	03027.jpg	04509.jpg	05964.jpg	07369.jpg	08778.jpg
00098.jpg	01598.jpg	03032.jpg	04512.jpg	05966.jpg	07373.jpg	08788.jpg
00099.jpg	01602.jpg	03033.jpg	04517.jpg	05967.jpg	07379.jpg	08791.jpg
00100.jpg	01606.jpg	03035.jpg	04523.jpg	05968.jpg	07380.jpg	08793.jpg
00105.jpg	01607.jpg	03036.jpg	04524.jpg	05969.jpg	07382.jpg	08795.jpg
00108.jpg	01608.jpg	03037.jpg	04527.jpg	05974.jpg	07385.jpg	08796.jpg
00109.jpg	01610.jpg	03040.jpg	04529.jpg	05977.jpg	07386.jpg	08797.jpg
00118.jpg	01615.jpg	03054.jpg	04530.jpg	05978.jpg	07387.jpg	08799.jpg
00119.jpg	01620.jpg	03055.jpg	04535.jpg	05979.jpg	07391.jpg	08805.jpg
00126.jpg	01626.jpg	03058.jpg	04536.jpg	05980.jpg	07392.jpg	08810.jpg
00127.jpg	01627.jpg	03063.jpg	04538.jpg	05981.jpg	07395.jpg	08811.jpg
00130.jpg	01631.jpg	03068.jpg	04540.jpg	05983.jpg	07397.jpg	08812.jpg
00133.jpg	01632.jpg	03071.jpg	04541.jpg	05984.jpg	07398.jpg	08816.jpg
00134.jpg	01638.jpg	03072.jpg	04543.jpg	05988.jpg	07399.jpg	08818.jpg
00135.jpg	01639.jpg	03073.jpg	04545.jpg	05989.jpg	07407.jpg	08828.jpg
00140.jpg	01640.jpg	03075.jpg	04546.jpg	05992.jpg	07409.jpg	08831.jpg
00141.jpg	01646.jpg	03076.jpg	04549.jpg	05993.jpg	07412.jpg	08832.jpg
00148.jpg	01650.jpg	03078.jpg	04550.jpg	05994.jpg	07415.jpg	08833.jpg
00152.jpg	01652.jpg	03080.jpg	04552.jpg	05998.jpg	07418.jpg	08836.jpg
00153.jpg	01654.jpg	03082.jpg	04553.jpg	05999.jpg	07423.jpg	08840.jpg
00157.jpg	01656.jpg	03084.jpg	04554.jpg	06000.jpg	07425.jpg	08841.jpg
00158.jpg	01660.jpg	03089.jpg	04556.jpg	06002.jpg	07426.jpg	08842.jpg
00159.jpg	01663.jpg	03092.jpg	04559.jpg	06003.jpg	07427.jpg	08846.jpg
00161.jpg	01665.jpg	03099.jpg	04562.jpg	06004.jpg	07428.jpg	08849.jpg
00161.jpg	01665.jpg	03099.jpg	04562.jpg	06004.jpg	07428.jpg	08849.jpg

Now that training has finished, let's have a look at what it actually predicts in a test image

```
# pick one image from the test set
img, _ = dataset_val[5]
# put the model in evaluation mode
model.eval()
with torch.no_grad():
    prediction = model([img.to(device)])
```

Printing the prediction shows that we have a list of dictionaries. Each element of the list corresponds to a different image. As we have a single image, there is a single dictionary in the list. The dictionary contains the predictions for the image we passed. In this case, we can see that it contains boxes, prediction

```
[{'boxes': tensor([[ 847.1060,  238.2487, 1259.9624,  544.6272],
       [ 67.4519, 218.1165, 828.6486, 588.8081],
       [ 415.7662, 222.1267, 1248.2137, 561.9327],
       [ 158.3181, 204.8467, 793.3710, 589.0566],
       [ 16.0116, 308.3200, 484.0447, 584.0211]], device='cuda:0'),
 'labels': tensor([2, 2, 2, 3, 2], device='cuda:0'),
 'scores': tensor([0.9667, 0.6619, 0.1202, 0.1194, 0.0722], device='cuda:0')}]
```

▼ Visualization of the predictions

Let's inspect an image and visualize the top predicted bounding box.

```
from google.colab.patches import cv2_imshow
from PIL import Image
import matplotlib.pyplot as plt
import torch
import torchvision.transforms as T
import torchvision
import numpy as np
import json
import cv2
import random
import warnings
warnings.filterwarnings('ignore')

# set to evaluation mode
model.eval()

file = open('3classes/3classes.json', 'r')
data = json.loads(file.read())
CLASS_NAMES = []
for category in data['categories']:
    CLASS_NAMES.append(category['name'])
print(CLASS_NAMES)

device = torch.device('cuda') if torch.cuda.is_available() else torch.device('cpu')
model.to(device)

def get_prediction(img_path, confidence):
    """
    get_prediction
    parameters:
        - img_path - path of the input image
        - confidence - threshold to keep the prediction or not
    method:
        - Image is obtained from the image path
        - the image is converted to image tensor using PyTorch's Transforms
        - image is passed through the model to get the predictions
    """
    pass
```

- classes and bounding boxes are obtained from the model and soft masks are made ie: eg. segment of cat is made 1 and rest of the image is made 0

"""

```
img = Image.open(img_path)
transform = T.Compose([T.ToTensor()])
img = transform(img)

img = img.to(device)
pred = model([img])
pred_score = list(pred[0]['scores'].detach().cpu().numpy())
pred_t = [pred_score.index(x) for x in pred_score if x>confidence][-1]
pred_class = [CLASS_NAMES[i] for i in list(pred[0]['labels'].cpu().numpy())]
pred_boxes = [[(i[0], i[1]), (i[2], i[3])] for i in list(pred[0]['boxes'].detach().cpu().numpy())]
pred_boxes = pred_boxes[:pred_t+1]
pred_class = pred_class[:pred_t+1]
return pred_boxes, pred_class
```

```
def detect_object(img_path, confidence=0.5, rect_th=2, text_size=1, text_th=2):
    """
```

detect_object

parameters:

- img_path - path to input image
 - confidence- confidence to keep the prediction or not
 - rect_th - rect thickness
 - text_size
 - text_th - text thickness
- method:
- prediction is obtained by get_prediction
 - each box is given a color
 - final output is displayed

"""

```
boxes, pred_cls = get_prediction(img_path, confidence)
```

```
img = cv2.imread(img_path)
```

```
img = cv2.cvtColor(img, cv2.COLOR_BGR2RGB)
```

```
for i in range(len(boxes)):
```

```
    random.seed(32)
```

```
    cv2.rectangle(img, boxes[i][0], boxes[i][1], color=(0,255,0), thickness=rect_th)
```

```
    cv2.putText(img,pred_cls[i], boxes[i][0], cv2.FONT_HERSHEY_SIMPLEX, text_size, (255,0,0), text_th)
```

```
return img
```

```
['None', 'dozer', 'dump_truck', 'excavator']
```

#Try with one image from the dataset

```
img_path = '3classes/10052.jpg'
```

```
img = detect_object(img_path, confidence=0.7)
```

```
img = cv2.cvtColor(img, cv2.COLOR_BGR2RGB)
```

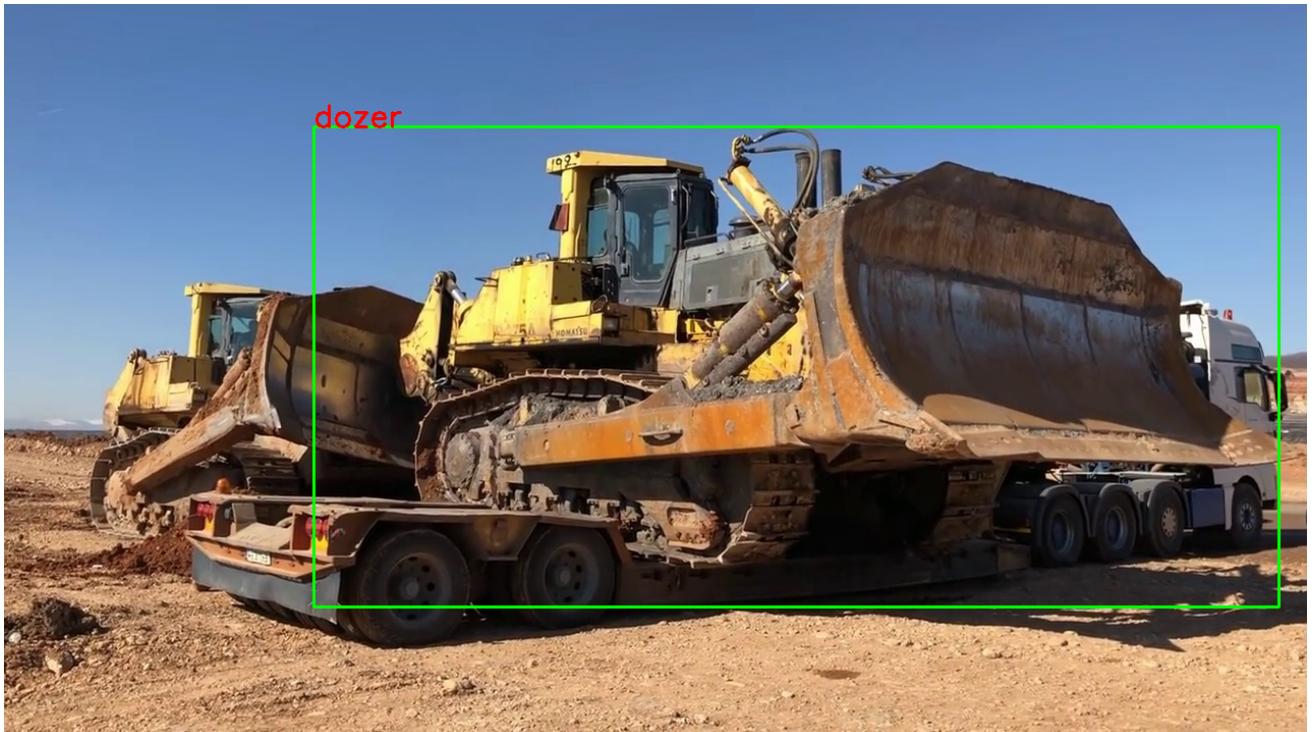
```
cv2_imshow(img)
```

```
img_path1 = '3classes/10071.jpg'
```

```
img1 = detect_object(img_path1, confidence=0.7)
```

```
img1 = cv2.cvtColor(img1, cv2.COLOR_BGR2RGB)
```

```
cv2_imshow(img1)
```



Looks pretty good!

Wrapping up

In this tutorial, you have learned how to create your own training pipeline for object detection models, on a custom dataset. For that, you wrote a `torch.utils.data.Dataset` class that returns the images and the ground truth boxes. You also leveraged a Faster R-CNN model pre-trained on COCO train2017 in order to perform transfer learning on this new dataset.

For a more complete example, which includes multi-machine / multi-gpu training, check `references/detection/train.py`, which is present in the [torchvision GitHub repo](#).

Reference: This handson was prepared following this [PyTorch](#) tutorial

▼ Give your judgement

Now its time for you to evaluate the model performance and comment on the following scenarios

1. Do the following hyper parameters affect the accuracy, precision and recall? If so, then explain how.

- i. learning rates
- ii. Optimizer (Can try with SGD, Adam, RMSProp)
- iii. batch size
- iv. Number of epochs

2. Try the model with another backbone and comment on your observations.

▼ Calculate accuracy, and draw precision-recall Curve

```
from __future__ import division
import warnings
import numpy as np
import matplotlib.pyplot as plt

def precision(tp, fn, fp, tn):
    with np.errstate(divide='ignore', invalid='ignore'):
        return tp/(tp + fp)

def recall(tp, fn, fp, tn):
    with np.errstate(divide='ignore', invalid='ignore'):
        return tp/(tp + fn)
```

```

def precision_gain(tp, fn, fp, tn):
    """Calculates Precision Gain from the contingency table
    This function calculates Precision Gain from the entries of the contingency
    table: number of true positives (TP), false negatives (FN), false positives
    (FP), and true negatives (TN). More information on Precision-Recall-Gain
    curves and how to cite this work is available at
    http://www.cs.bris.ac.uk/~flach/PRGcurves/.
    """
    n_pos = tp + fn
    n_neg = fp + tn
    with np.errstate(divide='ignore', invalid='ignore'):
        prec_gain = 1. - (n_pos/n_neg) * (fp/tp)
    if np.alen(prec_gain) > 1:
        prec_gain[tn + fn == 0] = 0
    elif tn + fn == 0:
        prec_gain = 0
    return prec_gain

def recall_gain(tp, fn, fp, tn):
    """Calculates Recall Gain from the contingency table
    This function calculates Recall Gain from the entries of the contingency
    table: number of true positives (TP), false negatives (FN), false positives
    (FP), and true negatives (TN). More information on Precision-Recall-Gain
    curves and how to cite this work is available at
    http://www.cs.bris.ac.uk/~flach/PRGcurves/.
    Args:
        tp (float) or ([float]): True Positives
        fn (float) or ([float]): False Negatives
        fp (float) or ([float]): False Positives
        tn (float) or ([float]): True Negatives
    Returns:
        (float) or ([float])
    """
    n_pos = tp + fn
    n_neg = fp + tn
    with np.errstate(divide='ignore', invalid='ignore'):
        rg = 1. - (n_pos/n_neg) * (fn/tp)
    if np.alen(rg) > 1:
        rg[tn + fn == 0] = 1
    elif tn + fn == 0:
        rg = 1
    return rg

def create_segments(labels, pos_scores, neg_scores):
    n = np.alen(labels)
    # reorder labels and pos_scores by decreasing pos_scores, using increasing neg_scores
    new_order = np.lexsort((neg_scores, -pos_scores))
    labels = labels[new_order]
    pos_scores = pos_scores[new_order]
    neg_scores = neg_scores[new_order]
    # create a table of segments
    segments = {'pos_score': np.zeros(n), 'neg_score': np.zeros(n),
                'pos_count': np.zeros(n), 'neg_count': np.zeros(n)}
    j = -1
    for i, label in enumerate(labels):
        if ((i == 0) or (pos_scores[i-1] != pos_scores[i]))

```

```

        or (neg_scores[i-1] != neg_scores[i])):
            j += 1
            segments['pos_score'][j] = pos_scores[i]
            segments['neg_score'][j] = neg_scores[i]
        if label == 0:
            segments['neg_count'][j] += 1
        else:
            segments['pos_count'][j] += 1
    segments['pos_score'] = segments['pos_score'][0:j+1]
    segments['neg_score'] = segments['neg_score'][0:j+1]
    segments['pos_count'] = segments['pos_count'][0:j+1]
    segments['neg_count'] = segments['neg_count'][0:j+1]
    return segments

def get_point(points, index):
    keys = points.keys()
    point = np.zeros(np.alen(keys))
    key_indices = dict()
    for i, key in enumerate(keys):
        point[i] = points[key][index]
        key_indices[key] = i
    return [point, key_indices]

def insert_point(new_point, key_indices, points, precision_gain=0,
                recall_gain=0, is_crossing=0):
    for key in key_indices.keys():
        points[key] = np.insert(points[key], 0, new_point[key_indices[key]])
    points['precision_gain'][0] = precision_gain
    points['recall_gain'][0] = recall_gain
    points['is_crossing'][0] = is_crossing
    new_order = np.lexsort((-points['precision_gain'], points['recall_gain']))
    for key in points.keys():
        points[key] = points[key][new_order]
    return points

def _create_crossing_points(points, n_pos, n_neg):
    n = n_pos+n_neg
    points['is_crossing'] = np.zeros(np.alen(points['pos_score']))
    # introduce a crossing point at the crossing through the y-axis
    j = np.amin(np.where(points['recall_gain'] >= 0)[0])
    if points['recall_gain'][j] > 0: # otherwise there is a point on the boundary and
        [point_1, key_indices_1] = get_point(points, j)
        [point_2, key_indices_2] = get_point(points, j-1)
        delta = point_1 - point_2
        if delta[key_indices_1['TP']] > 0:
            alpha = (n_pos*n_pos/n - points['TP'][j-1]) / delta[key_indices_1['TP']]
        else:
            alpha = 0.5
        with warnings.catch_warnings():
            warnings.simplefilter("ignore")
            new_point = point_2 + alpha*delta
        new_prec_gain = precision_gain(new_point[key_indices_1['TP']], new_point[key_in
                                         new_point[key_indices_1['FP']], new_point[key_in
    points = insert_point(new_point, key_indices_1, points,

```

```

# now introduce crossing points at the crossings through the non-negative part of t
x = points['recall_gain']
y = points['precision_gain']
temp_y_0 = np.append(y, 0)
temp_0_y = np.append(0, y)
temp_1_x = np.append(1, x)
with np.errstate(invalid='ignore'):
    indices = np.where(np.logical_and((temp_y_0 * temp_0_y < 0), (temp_1_x >= 0)))[0]
    for i in indices:
        cross_x = x[i-1] + (-y[i-1]) / (y[i] - y[i-1]) * (x[i] - x[i-1])
        [point_1, key_indices_1] = get_point(points, i)
        [point_2, key_indices_2] = get_point(points, i-1)
        delta = point_1 - point_2
        if delta[key_indices_1['TP']] > 0:
            alpha = (n_pos * n_pos / (n - n_neg * cross_x) - points['TP'][i-1]) / delta
        else:
            alpha = (n_neg / n_pos * points['TP'][i-1] - points['FP'][i-1]) / delta[key_indices_2['TP']]
        with warnings.catch_warnings():
            warnings.simplefilter("ignore")
            new_point = point_2 + alpha*delta
        new_rec_gain = recall_gain(new_point[key_indices_1['TP']], new_point[key_indices_1['FN']],
                                    new_point[key_indices_1['FP']], new_point[key_indices_1['TN']],
                                    points=insert_point(new_point, key_indices_1, points,
                                                         recall_gain=new_rec_gain, is_crossing=1))
        i += 1
        indices += 1
        x = points['recall_gain']
        y = points['precision_gain']
        temp_y_0 = np.append(y, 0)
        temp_0_y = np.append(0, y)
        temp_1_x = np.append(1, x)
    return points

def create_prg_curve(labels, pos_scores, neg_scores=[]):
    """Precision-Recall-Gain curve
    This function creates the Precision-Recall-Gain curve from the vector of
    labels and vector of scores where higher score indicates a higher
    probability to be positive. More information on Precision-Recall-Gain
    curves and how to cite this work is available at
    http://www.cs.bris.ac.uk/~flach/PRGcurves/.
    """
    create_crossing_points = True # do it always because calc_auprg otherwise gives the wrong result
    if np.alen(neg_scores) == 0:
        neg_scores = -pos_scores
    n = np.alen(labels)
    n_pos = np.sum(labels)
    n_neg = n - n_pos
    # convert negative labels into 0s
    labels = 1 * (labels == 1)
    segments = create_segments(labels, pos_scores, neg_scores)
    # calculate recall gains and precision gains for all thresholds
    points = dict()
    points['pos_score'] = np.insert(segments['pos_score'], 0, np.inf)
    points['neg_score'] = np.insert(segments['neg_score'], 0, -np.inf)
    ...

```

```

points['TP'] = np.insert(np.cumsum(segments['pos_count']), 0, 0)
points['FP'] = np.insert(np.cumsum(segments['neg_count']), 0, 0)
points['FN'] = n_pos - points['TP']
points['TN'] = n_neg - points['FP']
points['precision'] = precision(points['TP'], points['FN'], points['FP'], points['TN'])
points['recall'] = recall(points['TP'], points['FN'], points['FP'], points['TN'])
points['precision_gain'] = precision_gain(points['TP'], points['FN'], points['FP'], points['TN'])
points['recall_gain'] = recall_gain(points['TP'], points['FN'], points['FP'], points['TN'])
if create_crossing_points == True:
    points = _create_crossing_points(points, n_pos, n_neg)
else:
    points['pos_score'] = points['pos_score'][1:]
    points['neg_score'] = points['neg_score'][1:]
    points['TP'] = points['TP'][1:]
    points['FP'] = points['FP'][1:]
    points['FN'] = points['FN'][1:]
    points['TN'] = points['TN'][1:]
    points['precision_gain'] = points['precision_gain'][1:]
    points['recall_gain'] = points['recall_gain'][1:]
with np.errstate(invalid='ignore'):
    points['in_unit_square'] = np.logical_and(points['recall_gain'] >= 0,
                                                points['precision_gain'] >= 0)
return points

```

```

def calc_auprg(prg_curve):
    """Calculate area under the Precision-Recall-Gain curve
    This function calculates the area under the Precision-Recall-Gain curve
    from the results of the function create_prg_curve. More information on
    Precision-Recall-Gain curves and how to cite this work is available at
    http://www.cs.bris.ac.uk/~flach/PRGcurves/.
    """
    area = 0
    recall_gain = prg_curve['recall_gain']
    precision_gain = prg_curve['precision_gain']
    for i in range(1, len(recall_gain)):
        if (not np.isnan(recall_gain[i-1])) and (recall_gain[i-1] >= 0):
            width = recall_gain[i] - recall_gain[i-1]
            height = (precision_gain[i] + precision_gain[i-1]) / 2
            area += width * height
    return(area)

```

```

# from
def convex_hull(points):
    """Computes the convex hull of a set of 2D points.
    Input: an iterable sequence of (x, y) pairs representing the points.
    Output: a list of vertices of the convex hull in counter-clockwise order,
            starting from the vertex with the lexicographically smallest coordinates.
    Implements Andrew's monotone chain algorithm. O(n log n) complexity.
    Source code from:
    https://en.wikibooks.org/wiki/Algorithm\_Implementation/Geometry/Convex\_hull/Monotonicity\_Stack
    """

```

```

# Sort the points lexicographically (tuples are compared lexicographically).
# Remove duplicates to detect the case we have just one unique point.
points = sorted(set(points))

# Boring case: no points or a single point, possibly repeated multiple times.

```

```

if len(points) <= 1:
    return points

# 2D cross product of OA and OB vectors, i.e. z-component of their 3D cross product
# Returns a positive value, if OAB makes a counter-clockwise turn,
# negative for clockwise turn, and zero if the points are collinear.
def cross(o, a, b):
    return (a[0] - o[0]) * (b[1] - o[1]) - (a[1] - o[1]) * (b[0] - o[0])

# Build upper hull
upper = []
for p in reversed(points):
    while len(upper) >= 2 and cross(upper[-2], upper[-1], p) <= 0:
        upper.pop()
    upper.append(p)

return upper

```

def plot_prg(prg_curve, show_convex_hull=True, show_f_calibrated_scores=False):

 """Plot the Precision-Recall-Gain curve
This function plots the Precision-Recall-Gain curve resulting from the
function create_prg_curve using ggplot. More information on
Precision-Recall-Gain curves and how to cite this work is available at
<http://www.cs.bris.ac.uk/~flach/PRGcurves/>.
@param prg_curve the data structure resulting from the function create_prg_curve
@param show_convex_hull whether to show the convex hull (default: TRUE)
@param show_f_calibrated_scores whether to show the F-calibrated scores (default:TRUE)
@return the ggplot object which can be plotted using print()
@details This function plots the Precision-Recall-Gain curve, indicating
 for each point whether it is a crossing-point or not (see help on
 create_prg_curve). By default, only the part of the curve
 within the unit square [0,1]x[0,1] is plotted.
@examples
labels = c(1,1,1,0,1,1,1,1,1,0,1,1,1,0,0,1,0,0,0,1,0,1)
scores = (25:1)/25
plot_prg(create_prg_curve(labels,scores))
"""
pg = prg_curve['precision_gain']
rg = prg_curve['recall_gain']

fig = plt.figure(figsize=(6,5))
plt.clf()
plt.axes(frameon=False)
ax = fig.gca()
ax.set_xticks(np.arange(0,1.25,0.25))
ax.set_yticks(np.arange(0,1.25,0.25))
ax.grid(b=True)
ax.set_xlim((-0.05,1.02))
ax.set_ylim((-0.05,1.02))
ax.set_aspect('equal')
Plot vertical and horizontal lines crossing the 0 axis
plt.axvline(x=0, ymin=-0.05, ymax=1, color='k')
plt.axhline(y=0, xmin=-0.05, xmax=1, color='k')
plt.axvline(x=1, ymin=0, ymax=1, color='k')
plt.axhline(y=1, xmin=0, xmax=1, color='k')
Plot cyan lines
indices = np.arange(np.argmax(prg_curve['in_unit_square']) - 1,
 np.argmax(prg_curve['in_unit_square']) + 1)